

Step 1: Requirements

Functional requirements:

- Recursively fetches links from a set of starting pages
- Scalable such that it can crawl entire web
- Used to fetch hundreds of millions of web document
- Modular way with expectation that new functionality will be added e.g., new document types

Design considerations:

- Assume only HTTP for now
- Expected number of pages to crawl is 15B


Step 2: Back of Envelope Calculations

Storage

- 15B/(4w*7d*86400s)=6200pages/s
- 15B*(100KB+500)=1.5PB
- 1.5PB/0.7=2.14PB (don't want to go above 70% full)


Step 3: High-Level Design

- Basic algorithm:
 - Pick URL from unvisited URL list
 - Determine IP address of host
 - Establish connection and download doc
 - Parse doc to look for new URLs
 - Add new URLs to list of unvisited URLs
 - Process downloaded doc e.g., store it, index it, etc
 - Repeat
- Usually use BFS
- Path-ascending crawling can find non-linked paths
- What makes this difficult problem:
 - Large volume of web pages
 - Rate of change on web pages
- High level design needs:
 - URL frontier to store list of URLs to download; prioritize which URLs to crawl first
 - HTML fetcher to retrieve web page from server
 - Extractor to extract links from HTML documents
 - Duplicate eliminator to make sure same content is not extracted twice
 - Datastore to store retrieved pages, URLs and other metadata



Step 6: Detailed Design

- Assume crawler running on one server, and crawling done by multiple threads:
 - Remove absolute URL from shared URL frontier
 - For now, we just support HTTP, but should be designed modularly so future protocols can be supported
 - Thread downloads document and places it into a Document Input Stream (DIS)
 - This enables other modules to re-read the document multiple times
 - Worker invokes dedupe test to see if doc has been seen before
 - If so, doc is not processed and thread moves to next URL from URL frontier
- Next, crawler needs to process downloaded doc
 - Each doc can have different MIME types e.g., HTML, image, video
 - Implement MIME type modularly so we can move to other types later
 - HTML processing module extract all links from page
 - Each link converted to URL and tested against user-supplied URL filter to see if it should be downloaded
 - If passes filter, worker sees if URL has already been seen
 - If not then add to URL frontier



- URL frontier
 - Distribute into multiple servers
 - Hash function maps each URL to a server that will be responsible for crawling it
 - Hash function to map each hostname to a thread number
 - Each thread has a FIFO queue. This prevents overloading target server by downloading too many pages at once.
- HTML fetcher
 - Download doc corresponding to a given URL
- Document input stream (DIS)
 - Cache document so it can be accessed by multiple servers
 - Cache small documents entirely in memory, larger docs written temporarily to disk
- Document dedupe test
 - Calculate a 64-bit checksum on every processed document and store in database
 - MD5 or SHA
- URL filters
 - Control set of URLs that are downloaded
- Domain name resolution
 - Web crawler must use DNS to map webserver hostname into an IP address
- URL dedupe test
 - Dedupe links extracted from a URL document